

DDL Commands:

```
CREATE TABLE UserProfile(  
  userFirstName VARCHAR(100) NOT NULL,  
  userLastName VARCHAR(100) NOT NULL,  
  destinationCity VARCHAR(100) NOT NULL,  
  email VARCHAR(100) NOT NULL,  
  password VARCHAR(100) NOT NULL,  
  PRIMARY KEY(email),  
  FOREIGN KEY(destinationCity) REFERENCES AirportData(airportCity)  
);
```

```
CREATE TABLE CountryData(  
  country VARCHAR(100) NOT NULL,  
  countryCode VARCHAR(3) NOT NULL,  
  population INT,  
  region VARCHAR(100),  
  PRIMARY KEY(country)  
);
```

```
CREATE TABLE AirportData(  
  country VARCHAR(100) NOT NULL,  
  airportCity VARCHAR(100) NOT NULL,  
  airportName VARCHAR(100) NOT NULL,  
  airportCode VARCHAR(3) NOT NULL,  
  PRIMARY KEY(airportCode),  
  FOREIGN KEY(country) REFERENCES CountryData(country),  
  UNIQUE(airportCity),  
  UNIQUE(airportName)  
);
```

```
CREATE TABLE CovidCases(  
  country VARCHAR(100) NOT NULL,  
  countryCode VARCHAR(3),  
  date TIMESTAMP NOT NULL,  
  newCaseNumber INT,  
  newDeathNumber INT,  
  PRIMARY KEY(date, country),  
  FOREIGN KEY(country) REFERENCES CountryData(country)  
);
```

```
CREATE TABLE Vaccination(  
  country VARCHAR(100) NOT NULL,  
  countryCode VARCHAR(3),
```

```

date TIMESTAMP NOT NULL,
dailyVaccinationNumber INT,
PRIMARY KEY(date, country),
FOREIGN KEY(country) REFERENCES CountryData(country)
);

```

```

CREATE TABLE Hospitalization(
country VARCHAR(100) NOT NULL,
countryCode VARCHAR(3),
date TIMESTAMP NOT NULL,
patientNumber INT,
PRIMARY KEY(date, country),
FOREIGN KEY(country) REFERENCES CountryData(country)
);

```

```

CREATE TABLE Testing(
country VARCHAR(100) NOT NULL,
countryCode VARCHAR(3),
date TIMESTAMP NOT NULL,
newTestNumber INT,
PRIMARY KEY(date, country),
FOREIGN KEY(country) REFERENCES CountryData(country)
);

```

```

CREATE TABLE Ratings(
airportName VARCHAR(100) NOT NULL,
email VARCHAR(100) NOT NULL,
rating INT,
review TEXT,
PRIMARY KEY(airportName, email),
FOREIGN KEY(email) REFERENCES UserProfile(email),
FOREIGN KEY(airportName) REFERENCES AirportData(airportName)
);

```

Proof of Database:

```

mysql> SELECT table_name, table_rows FROM INFORMATION_SCHEMA.TABLES WHERE TABLE_SCHEMA = 'AbDB';
+-----+-----+
| TABLE_NAME | TABLE_ROWS |
+-----+-----+
| AirportData |          50 |
| CountryData |         224 |
| CovidCases  |       158812 |
| Hospitalization |       30595 |
| Ratings     |            0 |
| Testing     |        61187 |
| UserProfile |            0 |
| Vaccination |       74315 |
+-----+-----+
8 rows in set (0.02 sec)

```

Advanced SQL Commands:

```
mysql> SELECT country, SUM(newCaseNumber)/population as rate FROM CountryData NATURAL JOIN CovidCases GROUP BY country ORDER BY rate DESC LIMIT 15;
```

country	rate
Bahrain	0.7844
Israel	0.6007
Iceland	0.5885
Andorra	0.5578
Denmark	0.5551
San Marino	0.5155
Cyprus	0.4998
Seychelles	0.4905
Maldives	0.4865
Aruba	0.4708
Slovenia	0.4707
Netherlands	0.4681
Liechtenstein	0.4606
Switzerland	0.4492
Cayman Islands	0.4476

15 rows in set (0.26 sec)

```
mysql> SELECT airportName as 'Airport', country as 'Country', rate/3 AS 'Vaccination Rate' FROM (SELECT country as c, SUM(dailyVaccinationNumber)/population as rate FROM CountryData NATURAL JOIN Vaccination GROUP BY country) AS temp, AirportData WHERE rate > 0.5 AND country = c LIMIT 15;
```

Airport	Country	Vaccination Rate
Sydney Kingsford-Smith Airport	Australia	0.91413333
Toronto Pearson International Airport	Canada	0.82426667
Guangzhou Baiyun International Airport	China	0.81830000
Chongqing Jiangbei International Airport	China	0.81830000
Chengdu Shuangliu International Airport	China	0.81830000
Kunming Changshui International Airport	China	0.81830000
Beijing Capital International Airport	China	0.81830000
Shanghai Pudong International Airport	China	0.81830000
Shanghai Hongqiao International Airport	China	0.81830000
Shenzhen Bao'an International Airport	China	0.81830000
Xi'an Xianyang International Airport	China	0.81830000
Hong Kong International Airport	Hong Kong	0.68953333
Chhatrapati Shivaji Maharaj International Airport	India	0.55183333
Indira Gandhi International Airport	India	0.55183333
Soekarno-Hatta International Airport	Indonesia	0.49683333

15 rows in set (0.14 sec)

PRIMARY index:

PRIMARY and population index:

PRIMARY and newCaseNumber index:

[illegible]

PRIMARY, population, and newCaseNumber index:

```
| EXPLAIN
|
|-----+
|
|-----+
|
| -> Limit: 15 row(s) (actual time=280.504..280.506 rows=15 loops=1)
|   -> Sort: (sum(CovidCases.newCaseNumber) / CountryData.population) DESC, limit input to 15 row(s) per chunk (actual time=280.504..280.505 rows=15 loops=1)
|     -> Stream results (cost=54752.55 rows=15203) (actual time=2.369..280.164 rows=196 loops=1)
|       -> Group aggregate: sum(CovidCases.newCaseNumber) (cost=54752.55 rows=15203) (actual time=2.364..279.902 rows=196 loops=1)
|         -> Nested loop inner join (cost=53232.29 rows=15203) (actual time=0.357..255.278 rows=140504 loops=1)
|           -> Index scan on CountryData using PRIMARY (cost=23.15 rows=224) (actual time=0.046..0.200 rows=224 loops=1)
|           -> Filter: CovidCases.countryCode = CountryData.countryCode (cost=169.70 rows=68) (actual time=0.217..1.090 rows=627 loops=224)
|             -> Index lookup on CovidCases using country (country=CountryData.country) (cost=169.70 rows=679) (actual time=0.136..1.010 rows=689 loops=224)
|
|-----+
|
| 1 row in set (0.28 sec)
```

We tried two types of indexing, trying them separately and together. In short, both of them helped, but one more than the other

Our first index was on the population in the Country Data table. This affected only 224 rows however when we compared the result of this indexing to no indexing, our performance had increased by 6 seconds. Instead of taking 0.35 seconds, it took 0.29 seconds.

Our second index was on the newCaseNumber in the Covid Cases table - this is the largest table in our query, with approximately 158000 rows so we expected to see some results. We implemented this index after dropping our first index on population. Just by itself, it had increased our performance significantly, from 0.35 seconds to 0.28 seconds

We tried both the indices together and the result was actually the same as just the second index on newCaseNumber - 0.28 seconds.

Both of them helped increase our performance however newCaseNumber helped more, presumably because of how large the Covid Cases table is. Additionally, because we are doing a sum on the newCaseNumbers, it is fairly complex, meaning indexing would help greatly. The index on population also helped because our query was just that large.

Our query groups by country which is the primary key for CountryData. This means MySQL already created the PRIMARY index for it so our original query was already fast.

In conclusion, we will index on just newCaseNumber, as that provided the best results.

Second advanced SQL command:

PRIMARY index:

```
mysql> show index from CountryData;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| Table | Non_unique | Key_name | Seq_in_index | Column_name | Collation | Cardinality | Sub_part | Packed | Null | Index_type | Comment | Index_comment | Visible |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| CountryData | 0 | PRIMARY | 1 | country | A | 224 | NULL | NULL | NULL | BTREE | | | | YES
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

```
mysql> EXPLAIN ANALYZE SELECT airportName as 'Airport', country as 'Country', rate/3 AS 'Vaccination Rate' FROM (SELECT country as c, SUM(dailyVaccinationNumber)/population as rate FROM CountryData NATURAL JOIN Vaccination GROUP BY country) AS temp, AirportData WHERE rate > 0.5 AND country = c LIMIT 15;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| EXPLAIN
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| EXPLAIN
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-> Limit: 15 row(s) (cost=116.55 rows=15) (actual time=148.342..148.444 rows=15 loops=1)
-> Nested loop inner join (cost=116.55 rows=22) (actual time=148.341..148.442 rows=15 loops=1)
-> Table scan on temp (cost=0.12..2.77 rows=22) (actual time=0.002..0.008 rows=59 loops=1)
-> Materialize (cost=106.15..108.81 rows=22) (actual time=148.316..148.326 rows=163 loops=1)
-> Filter: ((sum(Vaccination.dailyVaccinationNumber) / CountryData.population) > 0.5) (cost=103.79 rows=22) (actual time=1.893..148.074 rows=163 loops=1)
-> Group aggregate: sum(Vaccination.dailyVaccinationNumber) (cost=103.79 rows=22) (actual time=1.075..147.902 rows=199 loops=1)
-> Nested loop inner join (cost=101.55 rows=22) (actual time=0.099..134.311 rows=76149 loops=1)
-> Index scan on CountryData using PRIMARY (cost=23.15 rows=224) (actual time=0.043..0.165 rows=224 loops=1)
-> Filter: (Vaccination.countryCode = CountryData.countryCode) (cost=0.25 rows=0) (actual time=0.060..0.573 rows=340 loops=224)
-> Index lookup on Vaccination using country (country=CountryData.country) (cost=0.25 rows=1) (actual time=0.016..0.530 rows=371 loops=224)
-> Index lookup on AirportData using country (country=temp.c) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=0 loops=59)
```

```
1 row in set (0.15 sec)
```

PRIMARY and population index:

[illegible]

```
mysql> CREATE INDEX dailyvacc ON Vaccination(dailyVaccinationNumber);
Query OK, 0 rows affected (0.50 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> SHOW INDEX FROM Vaccination;
```

Table	Non unique	Key name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment
Vaccination	0	PRIMARY	1	date	A	473	NULL	NULL		BTREE		
Vaccination	0	PRIMARY	2	country	A	74316	NULL	NULL		BTREE		
Vaccination	1	country	1	country	A	227	NULL	NULL		BTREE		
Vaccination	1	dailyvacc	1	dailyVaccinationNumber	A	38359	NULL	NULL	YES	BTREE		

```
4 rows in set (0.01 sec)
```

[illegible]

[illegible]

Our first index was on the population in the Country Data table. This affected only 224 rows so when we compared the result of this indexing to no indexing, our performance had actually decreased! Instead of taking 0.15 seconds, it took 0.16 seconds.

We tried both the indices together and the result actually took the longest - 0.18 seconds! This was by far the most interesting observation.

Our query groups by country which is the primary key for CountryData. This means MySQL already created the PRIMARY index for it so our original query was already very fast.

In conclusion, we will only index on the PRIMARY key, which is what we had originally, as it is either faster or the same as our other indices.